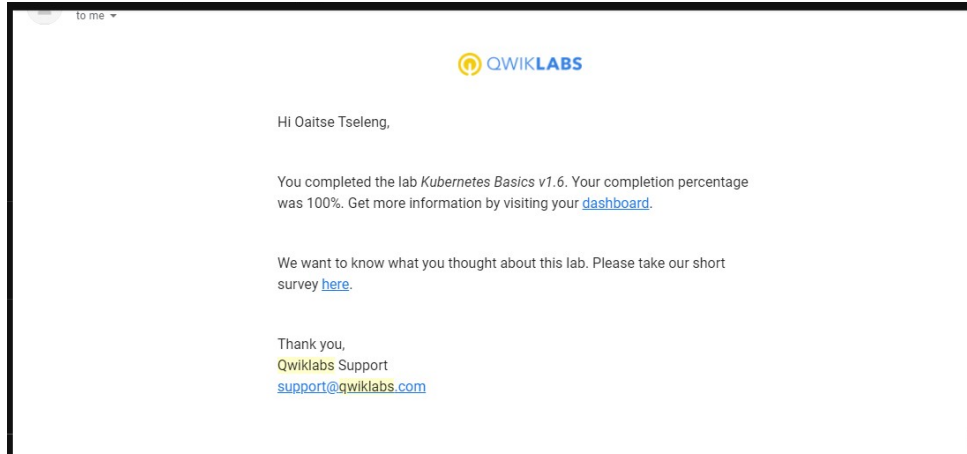


Part 1: Qwiklabs Completion(11 LABs)

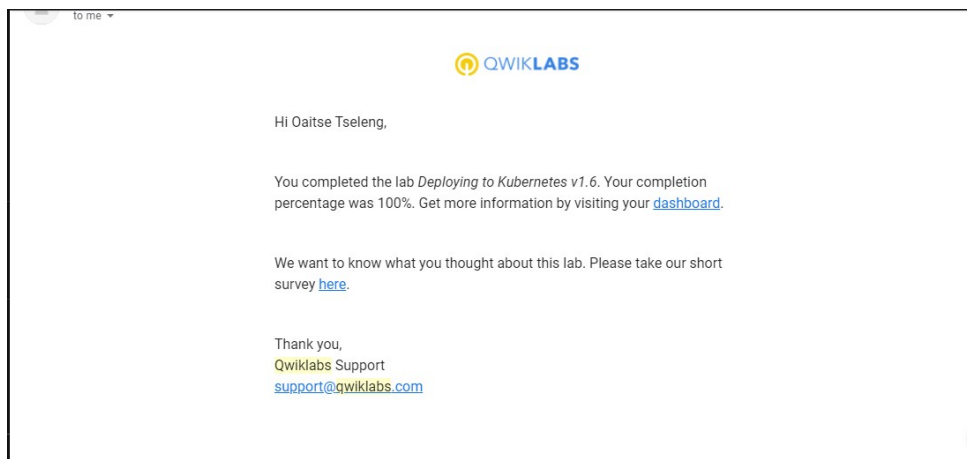
28 August:

1. Kubernetes Basics v1.6



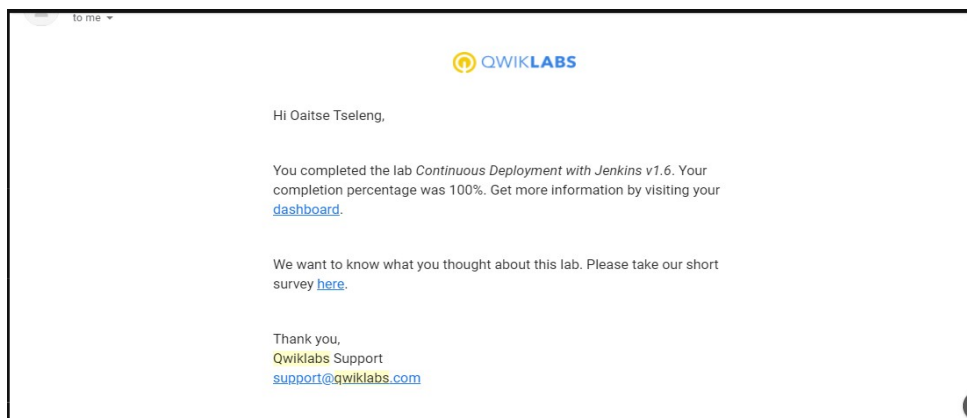
29 August:

2. Deploying to Kubernetes v1.6



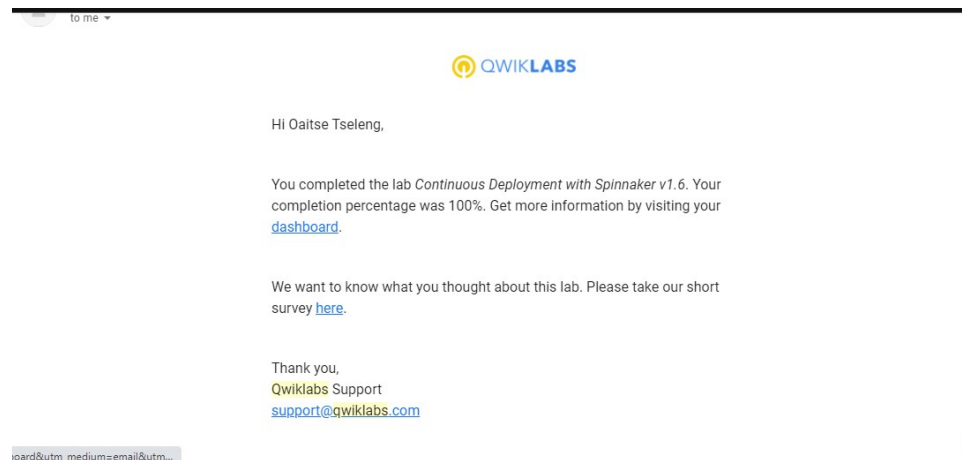
6 September:

3. Continuous Deployment with Jenkins v1.6

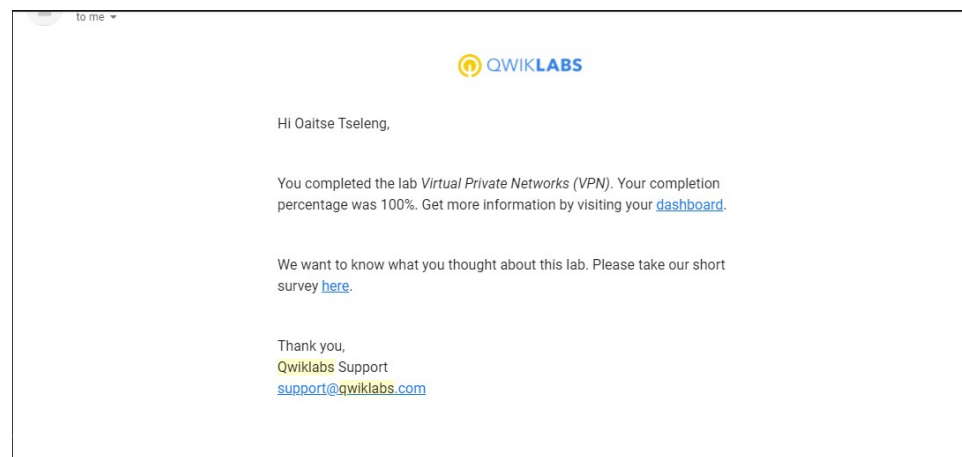


7 September:

4. Continuous Deployment with Spinnaker v1.6

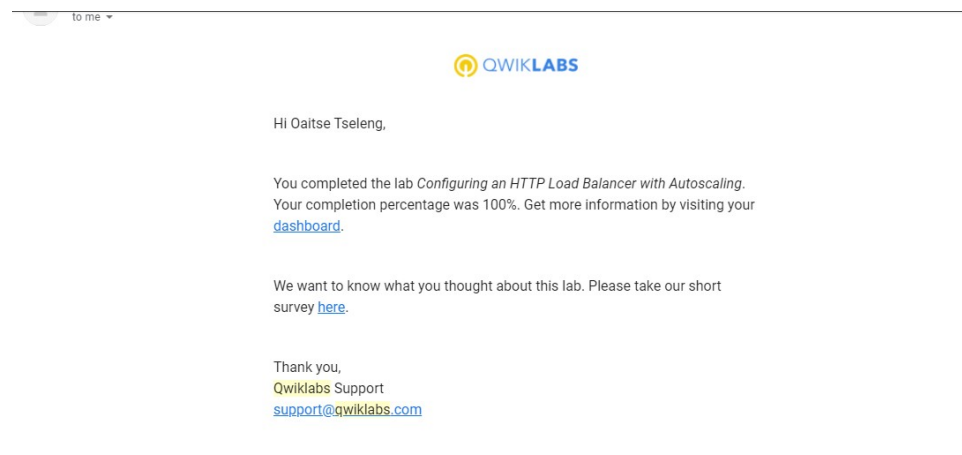


5. Virtual Private Networks (VPN)

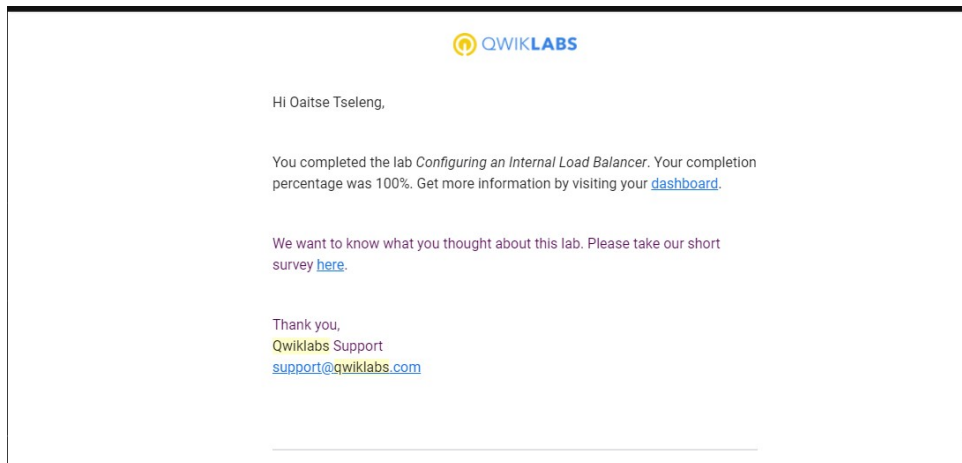


9 September:

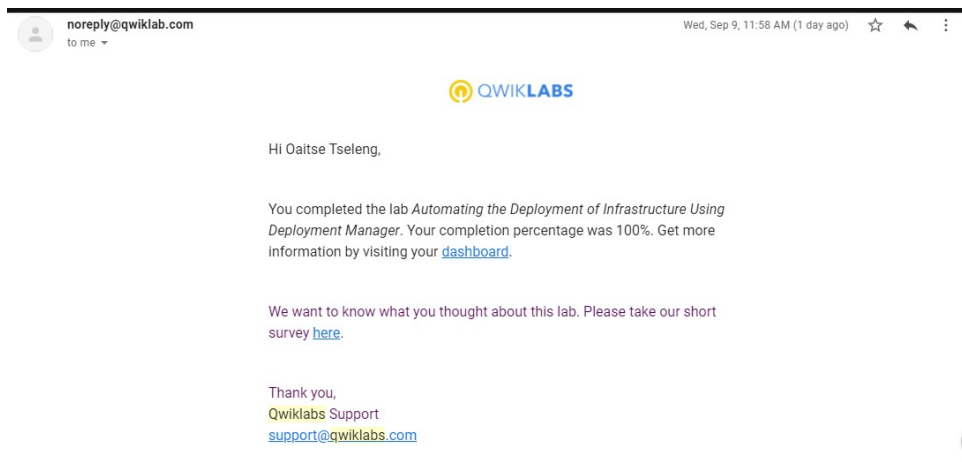
6. HTTP Load Balancer with Autoscaling



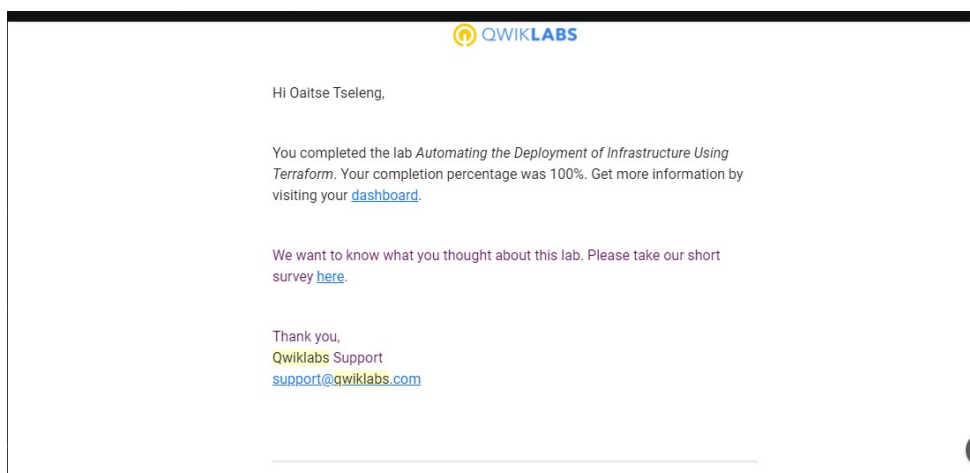
7. Configuring an Internal Load Balancer



8. Automating the Deployment of Infrastructure Using Deployment Manager

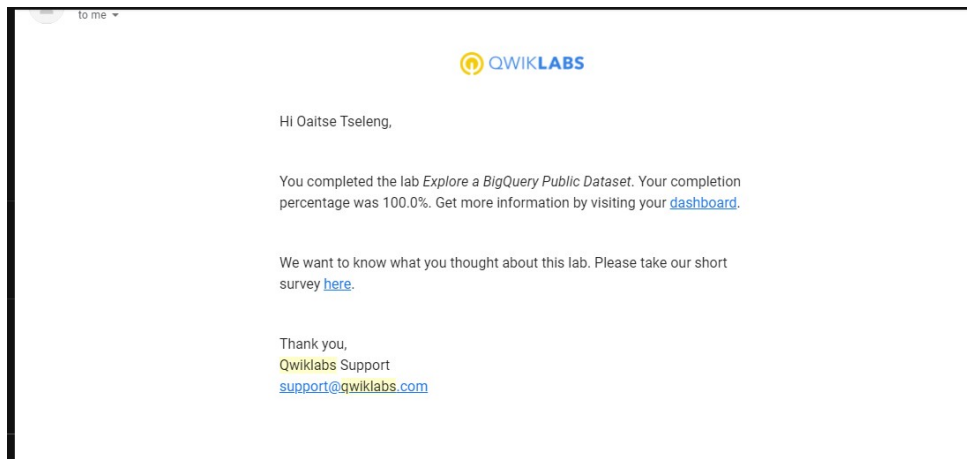


9. Automating the Deployment of Infrastructure Using Terraform



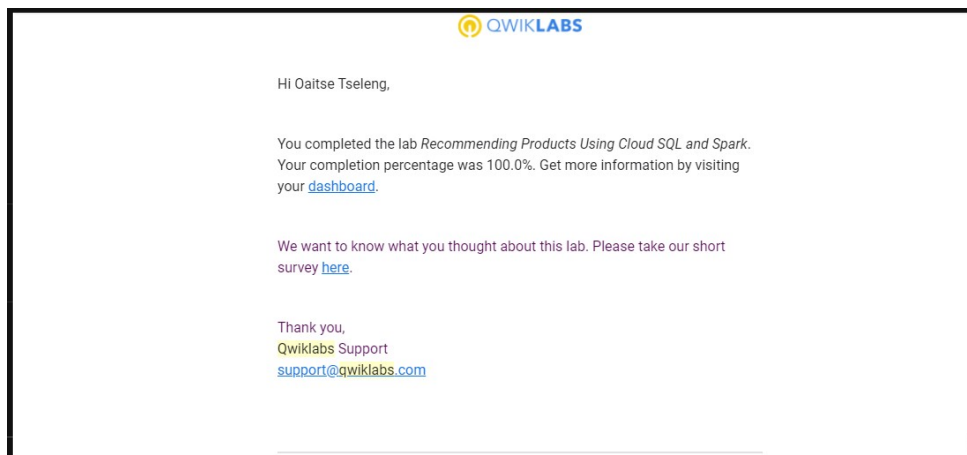
10 September:

10. BigQuery Public Dataset



11 September:

11. Recommending Products Using Cloud SQL and Spark



Phase 2: From Console to Command Line

Lab: Configuring an Internal Load Balancer

Task 1. Configure a health check firewall rule

- `gcloud compute firewall-rules create fw-allow-health-checks \ --network default \ --target-tags allow-health-checks \ --source-ranges 130.211.0.0/22, 35.191.0.0/16 \ --rules tcp:80`

Task 2: Create a NAT configuration using Cloud Router

- `gcloud compute routers nats create nat-config \ --router nat-router-us-central1`

Task 3: Create a custom image for a web server

- create vm instance:
`gcloud beta compute --project=qwiklabs-gcp-f72529adb527ca74 instances create webserver --zone=us-central1-a --machine-type=f1-micro --subnet=default --no-address --maintenance-policy=MIGRATE --service-account=984833531377-compute@developer.gserviceaccount.com --scopes=https://www.googleapis.com/auth/devstorage.read_only,https://www.googleapis.com/auth/logging.write,https://www.googleapis.com/auth/monitoring.write,https://www.googleapis.com/auth/servicecontrol,https://www.googleapis.com/auth/service.management.readonly,https://www.googleapis.com/auth/trace.append --tags=allow-health-checks --image=debian-10-buster-v20200910 --image-project=debian-cloud --boot-disk-size=10GB --no-boot-disk-auto-delete --boot-disk-type=pd-standard --boot-disk-device-name=webserver --no-shielded-secure-boot --no-shielded-vtpm --no-shielded-integrity-monitoring --reservation-affinity=any`
- `gcloud compute ssh webserver`
- Set the Apache service to start at boot:
`sudo apt-get update`
`sudo apt-get install -y apache2`
`sudo service apache2 start`
`curl localhost`
`sudo update-rc.d apache2 enable`
`sudo service apache2 status`
- Prepare the disk to create a custom image:
`gcloud compute instances describe webserver`
`gcloud compute instances delete webserver --zone us-central1-a`
`gcloud compute images create mywebserver \ --source-image webserver`

Task 4. Configure an instance template and create instance groups

- Configure instance template:
`gcloud beta compute --project=qwiklabs-gcp-f72529adb527ca74 instance-templates create mywebserver-template --machine-type=f1-micro --network=projects/qwiklabs-gcp-f72529adb527ca74/global/networks/default --no-address --maintenance-policy=MIGRATE --service-account=984833531377-compute@developer.gserviceaccount.com --scopes=https://www.googleapis.com/auth/devstorage.read_only,https://www.googleapis.com/auth/logging.write,https://www.googleapis.com/auth/monitoring.write,https://www.googleapis.com/auth/servicecontrol,https://`

www.googleapis.com/auth/service.management.readonly,<https://www.googleapis.com/auth/trace.append> --tags=allow-health-checks --image=mywebserver --image-project=qwiklabs-gcp-f72529adb527ca74 --boot-disk-size=10GB --boot-disk-type=pd-standard --boot-disk-device-name=mywebserver-template --no-shielded-secure-boot --shielded-vtpm --shielded-integrity-monitoring --reservation-affinity=any

- Create the managed instance groups:

us-central1:

```
gcloud compute --project "qwiklabs-gcp-f72529adb527ca74" health-checks create tcp "http-health-check" --timeout "5" --check-interval "10" --unhealthy-threshold "3" --healthy-threshold "2" --port "80"
```

```
gcloud beta compute --project=qwiklabs-gcp-f72529adb527ca74 instance-groups managed create us-central1-mig --base-instance-name=us-central1-mig --template=mywebserver-template --size=1 --zones=us-central1-b,us-central1-c,us-central1-f --instance-redistribution-type=PROACTIVE --health-check=http-health-check --initial-delay=60
```

```
gcloud beta compute --project "qwiklabs-gcp-f72529adb527ca74" instance-groups managed set-autoscaling "us-central1-mig" --region "us-central1" --cool-down-period "60" --max-num-replicas "2" --min-num-replicas "1" --target-load-balancing-utilization "0.8" --mode "on"
```

europe-west1:

```
gcloud beta compute --project=qwiklabs-gcp-f72529adb527ca74 instance-groups managed create europe-west1-mig --base-instance-name=europe-west1-mig --template=mywebserver-template --size=1 --zones=europe-west1-b,europe-west1-c,europe-west1-d --instance-redistribution-type=PROACTIVE --health-check=http-health-check --initial-delay=60
```

```
gcloud beta compute --project "qwiklabs-gcp-f72529adb527ca74" instance-groups managed set-autoscaling "europe-west1-mig" --region "europe-west1" --cool-down-period "60" --max-num-replicas "2" --min-num-replicas "1" --target-load-balancing-utilization "0.8" --mode "on"
```

Task 5. Configure the HTTP load balancer

- Create static ip addresses:
gcloud compute addresses create lb-ipv4-1 --ip-version=IPV4 --global
gcloud compute addresses create lb-ipv6-1 --ip-version=IPV6 --global
- Health checks:
gcloud compute health-checks create http http-health-check --port 80
- Create backend service:
gcloud compute backend-services create http-lb --protocol=HTTP --port-name=http --health-checks= http-health-check --global
- Add backends:
gcloud compute backend-services add-backend web-backend-service --instance-group=europe-west1-mig
gcloud compute backend-services add-backend web-backend-service --instance-group=us-central1-mig
gcloud compute forwarding-rules create http-content-rule --address=lb-ipv4-1 --global --target-http-proxy=http-lb-proxy --ports=80

- Stress test the HTTP load balancer

Creating new vm instance:

```
gcloud beta compute --project=qwiklabs-gcp-f72529adb527ca74 instances create stress-test
--zone=us-west1-c --machine-type=f1-micro --subnet=default --network-tier=PREMIUM --
maintenance-policy=MIGRATE --service-account=984833531377-
compute@developer.gserviceaccount.com
--scopes=https://www.googleapis.com/auth/devstorage.read_only,https://
www.googleapis.com/auth/logging.write,https://www.googleapis.com/auth/
monitoring.write,https://www.googleapis.com/auth/servicecontrol,https://
www.googleapis.com/auth/service.management.readonly,https://www.googleapis.com/
auth/trace.append --image=mywebserver --image-project=qwiklabs-gcp-f72529adb527ca74
--boot-disk-size=10GB --boot-disk-type=pd-standard --boot-disk-device-name=stress-test --
no-shielded-secure-boot --shielded-vtpm --shielded-integrity-monitoring --reservation-
affinity=any
```

create an environment variable for your load balancer IP address:

```
export LB_IP=<Enter your [LB_IP_v4] here>
```

place a load on the load balancer:

```
ab -n 500000 -c 1000 http://$LB_IP/
```

Lab: Exploring a Big Query Public Dataset

Task 1: Query a public dataset

- `bq query --use_legacy_sql=false \ "SELECT name, gender FROM bigquery-public-data.usa_names.usa_1910_2013 GROUP BY name, gender ORDER BY total DESC LIMIT 10"`

Task 3: Create Custom table

- `bq mk babynames`
- `bq ls`

Task 4: Load Data into new table

- `bq load / --source_format=CSV / babynames.names2014 / C:\Users\oait\Downloads\Compressed\yob2014.txt /name:string, gender:string, count:integer`

Task 5: Query the Table

- `bq query "SELECT name, count FROM `babynames.names_2014` WHERE gender = 'M' ORDER BY count DESC LIMIT 5"`

Lab: Recommend Products using ML with Cloud SQL and Dataproc

Task 1: Create a cloud SQL instance

- `gcloud sql instances create rentals --root-password=root`

Task 2: Create Tables

- gcloud sql connect rentals --user=root
- provide password
- run command "SHOW DATABASES;"
- Run Script
CREATE DATABASE IF NOT EXISTS recommendation_spark;

USE recommendation_spark;

DROP TABLE IF EXISTS Recommendation;
DROP TABLE IF EXISTS Rating;
DROP TABLE IF EXISTS Accommodation;

CREATE TABLE IF NOT EXISTS Accommodation
(
 id varchar(255),
 title varchar(255),
 location varchar(255),
 price int,
 rooms int,
 rating float,
 type varchar(255),
 PRIMARY KEY (ID)
);

CREATE TABLE IF NOT EXISTS Rating
(
 userId varchar(255),
 accId varchar(255),
 rating int,
 PRIMARY KEY(accId, userId),
 FOREIGN KEY (accId)
 REFERENCES Accommodation(id)
);

CREATE TABLE IF NOT EXISTS Recommendation
(
 userId varchar(255),
 accId varchar(255),
 prediction float,
 PRIMARY KEY(userId, accId),
 FOREIGN KEY (accId)
 REFERENCES Accommodation(id)
);

- SHOW DATABASES;
- USE recommendation_spark;
SHOW TABLES;
 - SELECT * FROM Accommodation;

Task 3: Stage data in cloud storage

- Creating bucket:
`gsutil mb gs://$DEVSHHELL_PROJECT_ID`
- Copying data from public dataset into our bucket:
`gsutil cp gs://cloud-training/bdml/v2.0/data/accommodation.csv`
`gsutil cp gs://cloud-training/bdml/v2.0/data/rating.csv`
- Show the files in our bucket:
`gsutil ls gs://$DEVSHHELL_PROJECT_ID`
- View some sample data:
`gsutil cat gs://$DEVSHHELL_PROJECT_ID/accommodation.csv`

Task 4: Load data from Cloud Storage into Cloud SQL tables

- `gcloud sql import csv rentals gs:// qwiklabs-gcp-02-d4e65ce05de3/accommodation.csv \ --database= recommendation_spark --table= Accommodation`
- `gcloud sql import csv rentals gs:// qwiklabs-gcp-02-d4e65ce05de3/ratings.csv \ --database= recommendation_spark --table= Ratings`

Task 5: Explore Cloud SQL data

- `USE recommendation_spark;`
- Query the ratings data:
`SELECT * FROM Rating LIMIT 15;`
- Count Number of rows in table:
`SELECT COUNT(*) AS num_ratings FROM Rating;`
- Average review rating of accommodations:
`SELECT`
`COUNT(userId) AS num_ratings,`
`COUNT(DISTINCT userId) AS distinct_user_ratings,`
`MIN(rating) AS worst_rating,`
`MAX(rating) AS best_rating,`
`AVG(rating) AS avg_rating`
`FROM Rating;`
- Users who have provided the most ratings:
`SELECT`
`userId,`
`COUNT(rating) AS num_ratings`
`FROM Rating`
`GROUP BY userId`
`ORDER BY num_ratings DESC;`

Task 6. Launch Dataproc

- Create dataproc cluster:
`gcloud dataproc clusters create rentals --region=us-central1 --master-machine-type= n1-standard-2 --worker-machine-type= n1-standard-2`
- Authorising cloud dataproc to connect to sql

```
echo "Authorizing Cloud Dataproc to connect with Cloud SQL"
CLUSTER=rentals
```

```

CLOUDSQL=rentals
ZONE=us-central1-a
NWORKERS=2

machines="$CLUSTER-m"
for w in `seq 0 $((NWORKERS - 1))`; do
    machines="$machines $CLUSTER-w-$w"
done

echo "Machines to authorize: $machines in $ZONE ... finding their IP addresses"
ips=""
for machine in $machines; do
    IP_ADDRESS=$(gcloud compute instances describe $machine --zone=$ZONE --
format='value(networkInterfaces.accessConfigs[].natIP)' | sed "s/\[//g" | sed "s/\]/g" )/32
    echo "IP address of $machine is $IP_ADDRESS"
    if [ -z $ips ]; then
        ips=$IP_ADDRESS
    else
        ips="$ips,$IP_ADDRESS"
    fi
done

echo "Authorizing [$ips] to access cloudsql=$CLOUDSQL"
gcloud sql instances patch $CLOUDSQL --authorized-networks $ips

```

Task 7. Run the ML model

- `gsutil cp gs://cloud-training/bdml/v2.0/model/train_and_apply.py train_and_apply.py`
`cloudshell edit train_and_apply.py`
- Edit apply and train file:
MAKE EDITS HERE
`CLOUDSQL_INSTANCE_IP = '<paste-your-cloud-sql-ip-here>' # <---- CHANGE (database server IP)`
`CLOUDSQL_PWD = '<type-your-cloud-sql-password-here>' # <---- CHANGE`
- Copy file to cloud stoage:
`gsutil cp train_and_apply.py gs://$DEVSHHELL_PROJECT_ID`

Task 8. Run your ML job on Dataproc

- `gcloud dataproc jobs submit pyspark gs:// qwiklabs-gcp-02-d4e65ce05de3 /train_and_apply.py --cluster=rentals`

Task 9. Explore inserted rows with SQL

- `USE recommendation_spark;`
- Check to see if dataproc finished its job:
`SELECT COUNT(*) AS count FROM Recommendation;`
- Find the recommendations for a user:
`SELECT r.userid, r.accoid, r.prediction, a.title, a.location, a.price, a.rooms, a.rating, a.type`

```
FROM Recommendation as r JOIN Accommodation as a  
ON r.accoid = a.id WHERE r.userid = 10;
```